

HW 5

Enter your name and EID here: Joseph Hendrix | jlh7459

You will submit this homework assignment as a pdf file on Gradescope.

For all questions, include the R commands/functions that you used to find your answer (show R chunk). Answers without supporting code will not receive credit. Write full sentences to describe your findings.

Question 1: (1 pt)

The dataset `world_bank_pop` is a built-in dataset in `tidyverse`. It contains information about total population and population growth, overall and more specifically in urban areas, for countries around the world. Take a look at it with `head()`. Is the data tidy? Why or why not?

```
# Call tidyr, dplyr and ggplot2 packages within tidyverse
library(tidyverse)
```

```
# Take a Look!
head(world_bank_pop)
```

```
## # A tibble: 6 × 20
##   country indica...1 `2000` `2001` `2002` `2003` `2004` `2005` `2006` `2007`
##   <chr>    <chr>    <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 ABW     SP.URB.... 4.24e4 4.30e4 4.37e4 4.42e4 4.47e+4 4.49e+4 4.49e+4 4.47e+4
## 2 ABW     SP.URB.... 1.18e0 1.41e0 1.43e0 1.31e0 9.51e-1 4.91e-1 -1.78e-2 -4.35e-1
## 3 ABW     SP.POP.... 9.09e4 9.29e4 9.50e4 9.70e4 9.87e+4 1.00e+5 1.01e+5 1.01e+5
## 4 ABW     SP.POP.... 2.06e0 2.23e0 2.23e0 2.11e0 1.76e+0 1.30e+0 7.98e-1 3.84e-1
## 5 AFG     SP.URB.... 4.44e6 4.65e6 4.89e6 5.16e6 5.43e+6 5.69e+6 5.93e+6 6.15e+6
## 6 AFG     SP.URB.... 3.91e0 4.66e0 5.13e0 5.23e0 5.12e+0 4.77e+0 4.12e+0 3.65e+0
## # ... with 10 more variables: `2008` <dbl>, `2009` <dbl>, `2010` <dbl>,
## #   `2011` <dbl>, `2012` <dbl>, `2013` <dbl>, `2014` <dbl>, `2015` <dbl>,
## #   `2016` <dbl>, `2017` <dbl>, and abbreviated variable name 1indicator
```

The data is tidy; every variable has its own column, every observation has its own row, and every value has its own cell.

Question 2: (1 pt)

Using `dplyr` functions on `world_bank_pop`, count how many distinct countries there are in the dataset. Does this makes sense? Why or why not?

```
# pull distinct countries from dataset
world_bank_pop %>%
  summarize(n_distinct(country))
```

```
## # A tibble: 1 × 1
##   `n_distinct(country)`
##               <int>
## 1                 264
```

There are 264 distinct countries in the dataset. This doesn't make sense as there are only 189 members of the World Bank.

Question 3: (2 pts)

Use one of the `pivot` functions on `world_bank_pop` to create a new dataset with the years 2000 to 2017 appearing as a *numeric* variable `year`, and the different values for the indicator variable are in a variable called `value`. Save this new dataset in your environment as `myworld1`.

```
# pivot dataset to elongate year variable
myworld1 <- pivot_longer(world_bank_pop,
                          cols = c(3:20),
                          names_to = "year",
                          values_to = "value")
```

```
myworld1
```

```
## # A tibble: 19,008 × 4
##   country indicator  year  value
##   <chr>    <chr>    <chr> <dbl>
## 1 ABW      SP.URB.TOTL 2000  42444
## 2 ABW      SP.URB.TOTL 2001  43048
## 3 ABW      SP.URB.TOTL 2002  43670
## 4 ABW      SP.URB.TOTL 2003  44246
## 5 ABW      SP.URB.TOTL 2004  44669
## 6 ABW      SP.URB.TOTL 2005  44889
## 7 ABW      SP.URB.TOTL 2006  44881
## 8 ABW      SP.URB.TOTL 2007  44686
## 9 ABW      SP.URB.TOTL 2008  44375
## 10 ABW     SP.URB.TOTL 2009  44052
## # ... with 18,998 more rows
```

How many lines are there per country? Why does it make sense?

```
myworld1 %>%
  filter(country == "ABW") %>%
  summarize(n())
```

```
## # A tibble: 1 × 1
##   `n()`
##   <int>
## 1     72
```

There are 72 rows for each country. This makes sense! There's 18 total years and 4 indicators for each year, for a total of 72.

Question 4: (3 pts)

Use another pivot function on `myworld1` to create a new dataset, `myworld2`, with the different categories for the indicator variable appearing as their own variables. Use `dplyr` functions to rename `SP.POP.GROW` and `SP.URB.GROW`, as `pop_growth` and `pop_urb_growth` respectively.

```
# define codes for observation refactoring
codes = list("SP.URB.TOTL" = "pop_urb_total", "SP.URB.GROW" = "pop_urb_grow", "SP.POP.TOTL" = "pop_total", "SP.POP.GROW" = "pop_growth")


myworld2 <- myworld1 %>%
  mutate(indicator = recode(indicator, !!!codes)) %>%      # refactor indicators
  pivot_wider(
    names_from = indicator,
    values_from = value
  )

myworld2
```

```
## # A tibble: 4,752 × 6
##   country year  pop_urb_total pop_urb_grow pop_total pop_growth
##   <chr>   <chr>      <dbl>      <dbl>      <dbl>      <dbl>
## 1 ABW     2000         42444         1.18       90853       2.06
## 2 ABW     2001         43048         1.41       92898       2.23
## 3 ABW     2002         43670         1.43       94992       2.23
## 4 ABW     2003         44246         1.31       97017       2.11
## 5 ABW     2004         44669         0.951      98737       1.76
## 6 ABW     2005         44889         0.491     100031       1.30
## 7 ABW     2006         44881        -0.0178    100832       0.798
## 8 ABW     2007         44686        -0.435    101220       0.384
## 9 ABW     2008         44375        -0.698    101353       0.131
## 10 ABW    2009         44052        -0.731    101453       0.0986
## # ... with 4,742 more rows
```

Using `dplyr` functions, find the ratio of urban growth compared to the population growth in the world for each year. *Hint: the country code `WLD` represents the entire world.* Create a `ggplot` to display how the percentage of urban population growth has changed over the years. Why does your graph not contradict the fact that the urban population worldwide is increasing over the years?

```
myworld2 %>%
  filter(country == "WLD") %>%
  mutate(ratio = pop_urb_grow / pop_growth) %>%      # calculate growth ratio
  ggplot(aes(x = year, y = ratio)) +
  geom_col()
```



```
myworld2 %>%  
  filter(country == "WLD") %>%  
  mutate(ratio = pop_urb_grow / pop_growth)
```

```
## # A tibble: 18 × 7
##   country year  pop_urb_total pop_urb_grow  pop_total pop_growth ratio
##   <chr>   <chr>      <dbl>      <dbl>      <dbl>    <dbl> <dbl>
## 1 WLD     2000      2858130756      2.16  6121682736      1.32  1.63
## 2 WLD     2001      2923079567      2.27  6201340258      1.30  1.75
## 3 WLD     2002      2991628819      2.35  6280530065      1.28  1.84
## 4 WLD     2003      3061267131      2.33  6359899296      1.26  1.84
## 5 WLD     2004      3132261863      2.32  6439825381      1.26  1.85
## 6 WLD     2005      3204583153      2.31  6520298763      1.25  1.85
## 7 WLD     2006      3277558354      2.28  6601476541      1.25  1.83
## 8 WLD     2007      3351069648      2.24  6683223772      1.24  1.81
## 9 WLD     2008      3426964053      2.26  6766296679      1.24  1.82
## 10 WLD    2009      3503454963      2.23  6849569339      1.23  1.81
## 11 WLD    2010      3580569790      2.20  6932869743      1.22  1.81
## 12 WLD    2011      3655009162      2.08  7014983968      1.18  1.76
## 13 WLD    2012      3730938916      2.08  7099557649      1.21  1.72
## 14 WLD    2013      3808101888      2.07  7185137526      1.21  1.72
## 15 WLD    2014      3886498272      2.06  7271322821      1.20  1.72
## 16 WLD    2015      3966059373      2.05  7357559450      1.19  1.73
## 17 WLD    2016      4046606978      2.03  7444157356      1.18  1.73
## 18 WLD    2017      4127612962      2.00  7530360149      1.16  1.73
```

The data shows a pretty consistent ratio of urban growth to population growth. This is consistent with an increase in global urban population over the years.

Question 5: (1 pt)

In `myworld2`, which country code had the highest population growth in 2017?

```
# sort myworld2 by population growth from 2017
myworld2 %>%
  filter(year == 2017) %>%
  group_by(country) %>%
  arrange(desc(pop_growth))
```

```
## # A tibble: 264 × 6
## # Groups:   country [264]
##   country year  pop_urb_total pop_urb_grow pop_total pop_growth
##   <chr>   <chr>      <dbl>      <dbl>      <dbl>      <dbl>
## 1 OMN     2017      3874061      5.95     4636262      4.67
## 2 BHR     2017      1331176      4.73     1492584      4.62
## 3 NRU     2017       13649      4.50       13649      4.50
## 4 NER     2017     3511546      4.18    21477348      3.82
## 5 GNQ     2017       908248      4.42     1267689      3.71
## 6 AGO     2017     19311773      4.38    29784193      3.31
## 7 UGA     2017     9942492      5.76    42862958      3.26
## 8 COD     2017     35691987      4.57    81339988      3.25
## 9 BDI     2017     1380411      5.72    10864245      3.18
## 10 TZA    2017     18942681      5.28    57310019      3.08
## # ... with 254 more rows
```

Oman (OMN) has the highest population growth in 2017 at 4.669.

Question 6: (1 pt)

When answering the previous, we only reported the three-letter code and (probably) have no idea what the actual country is. We will now use the package `countrycode` with a built-in dataset called `odelist` that has information about the coding system used by the World bank:

```
# Paste and run the following into your console (NOT HERE): install.packages("countrycode")

# Call the countrycode package
library(countrycode)

# Create a list of codes with matching country names
mycodes <- oodelist # continue this code...
```

Using `dplyr` functions, modify `mycodes` above to only keep the variables `continent`, `wb` (World Bank code), and `country.name.en` (country name in English). Then remove countries with missing `wb` code.

How many countries are there in `mycodes` ?

```
# refactor mycodes to cull irrelevant info
mycodes = mycodes %>%
  select(continent, wb, country.name.en) %>%
  filter(wb != "")

mycodes
```

```
## # A tibble: 218 × 3
##   continent wb   country.name.en
##   <chr>      <chr> <chr>
## 1 Asia      AFG   Afghanistan
## 2 Europe    ALB   Albania
## 3 Africa    DZA   Algeria
## 4 Oceania   ASM   American Samoa
## 5 Europe    AND   Andorra
## 6 Africa    AGO   Angola
## 7 Americas  ATG   Antigua & Barbuda
## 8 Americas  ARG   Argentina
## 9 Asia      ARM   Armenia
## 10 Americas ABW   Aruba
## # ... with 208 more rows
```

There are 218 countries in mycodes.

Question 7: (1 pt)

Use a `left_join()` function to add the information of the country codes **to** `myworld2` dataset. Match the two datasets based on the World Bank code. *Note: the World Bank code does not have the same name in each dataset.* Using `dplyr` functions, only keep the data available for Europe and for the year 2017. Save this new dataset as `myeurope`.

```
# join both datasets
myeurope = myworld2 %>%
  filter(year == 2017) %>%      # only from 2017
  left_join(mycodes %>%
    filter(continent == "Europe") %>%      # only from Europe
    select(wb, country.name.en)), by = c("country" = "wb")) %>%      # join on country
name
  filter(country.name.en != "") %>%      # get rid of missing non-Europe values
  select(country.name.en, country, everything())      # reorder dataset

myeurope
```

```
## # A tibble: 46 × 7
##   country.name.en    country year  pop_urb_total pop_urb_grow pop_t...1 pop_g...2
##   <chr>            <chr>   <chr>      <dbl>      <dbl>    <dbl>    <dbl>
## 1 Albania          ALB     2017      1706345      1.54    2.87e6 -0.0920
## 2 Andorra          AND     2017       67845     -0.520   7.70e4 -0.410
## 3 Austria          AUT     2017     5117624      1.15    8.81e6  0.827
## 4 Belgium          BEL     2017    11140192      0.401   1.14e7  0.358
## 5 Bulgaria          BGR     2017     5283572     -0.273   7.08e6 -0.730
## 6 Bosnia & Herzegovina BIH     2017     1679019      0.472   3.51e6 -0.279
## 7 Belarus          BLR     2017     7428883      0.674   9.51e6  0.0667
## 8 Switzerland      CHE     2017     6244619      1.13    8.47e6  1.10
## 9 Czechia          CZE     2017     7803157      0.379   1.06e7  0.236
## 10 Germany          DEU     2017     63890984     0.468   8.27e7  0.420
## # ... with 36 more rows, and abbreviated variable names 1pop_total, 2pop_growth
```

How many rows are there in this new dataset `myeurope`? What does each row represent?

```
# pull row count from myeurope
myeurope %>%
  summarize(n())
```


```
## # A tibble: 1 × 1
##   `n()`
##   <int>
## 1     46
```

There are 46 rows in `myeurope`, each row representing a country.

Question 8: (2 pts)

Using `dplyr` functions on `myeurope`, only keep information for the population growth in 2017 then compare the population growth per country with `ggplot` using `geom_bar()`. Make sure to order countries in order of population growth. Which country in Europe had the lowest population growth in 2017?

```
# build population growth bar graph
myeurope %>%
  select(country.name.en, pop_growth) %>%
  arrange(pop_growth) %>%
  ggplot(aes(x = reorder(country.name.en, pop_growth), y = pop_growth)) +
  geom_bar(stat = "identity") +
  coord_flip()
```

Lithuania has the lowest population growth.

Question 9: (1 pt)

When dealing with location data, we can actually visualize information on a map if we have geographic information such as latitude and longitude. Next, we will use a built-in function called `map_data()` to get geographic coordinates about countries in the world (see below). Take a look at the dataset `mapWorld`. What variables could we use to join `mapWorld` and `myeurope`? *Note: the variables do not have the same name in each dataset but they contain the same information.*

```
# Geographic coordinates about countries in the world
mapWorld <- map_data("world")

mapWorld
```

```
##           long      lat group order      region subregion
## 1 -69.89912 12.45200      1      1      Aruba      <NA>
## 2 -69.89571 12.42300      1      2      Aruba      <NA>
## 3 -69.94219 12.43853      1      3      Aruba      <NA>
## 4 -70.00415 12.50049      1      4      Aruba      <NA>
## 5 -70.06612 12.54697      1      5      Aruba      <NA>
## 6 -70.05088 12.59707      1      6      Aruba      <NA>
## 7 -70.03511 12.61411      1      7      Aruba      <NA>
## 8 -69.97314 12.56763      1      8      Aruba      <NA>
## 9 -69.91181 12.48047      1      9      Aruba      <NA>
## 10 -69.89912 12.45200      1     10      Aruba      <NA>
## 12  74.89131 37.23164      2     12 Afghanistan <NA>
## 13  74.84023 37.22505      2     13 Afghanistan <NA>
## 14  74.76738 37.24917      2     14 Afghanistan <NA>
## 15  74.73896 37.28564      2     15 Afghanistan <NA>
## 16  74.72666 37.29072      2     16 Afghanistan <NA>
## 17  74.66895 37.26670      2     17 Afghanistan <NA>
## [ reached 'max' / getOption("max.print") -- omitted 99322 rows ]
```

myeurope

```
## # A tibble: 46 × 7
##   country.name.en      country year  pop_urb_total pop_urb_grow pop_t...1 pop_g...2
##   <chr>             <chr>   <chr>      <dbl>      <dbl>    <dbl>    <dbl>
## 1 Albania          ALB     2017      1706345      1.54    2.87e6 -0.0920
## 2 Andorra          AND     2017       67845     -0.520    7.70e4 -0.410
## 3 Austria          AUT     2017      5117624      1.15    8.81e6  0.827
## 4 Belgium          BEL     2017     11140192      0.401    1.14e7  0.358
## 5 Bulgaria          BGR     2017      5283572     -0.273    7.08e6 -0.730
## 6 Bosnia & Herzegovina BIH     2017      1679019      0.472    3.51e6 -0.279
## 7 Belarus          BLR     2017      7428883      0.674    9.51e6  0.0667
## 8 Switzerland      CHE     2017      6244619      1.13    8.47e6  1.10
## 9 Czechia          CZE     2017      7803157      0.379    1.06e7  0.236
## 10 Germany          DEU     2017      63890984      0.468    8.27e7  0.420
## # ... with 36 more rows, and abbreviated variable names 1pop_total, 2pop_growth
```

We could use region from mapWorld to join on country.name.en in myeurope.

Question 10: (2 pts)

Use a joining function to check if any information from myeurope is not contained in mapWorld, matching the two datasets based on the country name.

```
# check where both datasets don't match
myeurope %>%
  anti_join(mapWorld, by = c("country.name.en" = "region"))
```

```
## # A tibble: 4 × 7
##   country.name.en    country year  pop_urb_total pop_urb_grow pop_to...1 pop_g...2
##   <chr>            <chr>   <chr>      <dbl>      <dbl>    <dbl>    <dbl>
## 1 Bosnia & Herzegovina BIH     2017      1679019      0.472  3507017  -0.279
## 2 Czechia          CZE     2017      7803157      0.379 10591323   0.236
## 3 United Kingdom   GBR     2017     54892898      0.958 66022273   0.648
## 4 Gibraltar        GIB     2017       34571      0.473   34571    0.473
## # ... with abbreviated variable names 1pop_total, 2pop_growth
```

Some countries such as United Kingdom did not have a match. Why do you think this happened? *Hint: find the distinct country names in `mapWorld`, arrange them in alphabetical order, and scroll through the names. Can you find any of these countries with no match in a slightly different form?*

```
# sort mapWorld by alphabetical region
mapWorld %>%
  group_by(region) %>%
  summarize(mean(group))
```

```
## # A tibble: 252 × 2
##   region      `mean(group)`
##   <chr>          <dbl>
## 1 Afghanistan         2
## 2 Albania              6
## 3 Algeria            486
## 4 American Samoa      24
## 5 Andorra             10
## 6 Angola              3.12
## 7 Anguilla             5
## 8 Antarctica         100.
## 9 Antigua            137
## 10 Argentina          19.8
## # ... with 242 more rows
```

It looks like the data from `myeurope` is older than `mapWorld`. Czechia in `myeurope` is now called Czech Republic, which is how it shows up in `mapWorld`. There's also small syntactic differences, like an ampersand instead of the word 'and'.

Question 11: (1 pt)

Consider the `myeurope` dataset. Recode some of the country names so that the countries with no match from the previous question (with the exception of Gibraltar which is not technically a country anyway) will have a match.

Hint: use `recode()` inside `mutate()` as described in this article <https://www.statology.org/recode-dplyr/> (<https://www.statology.org/recode-dplyr/>). Then add a pipe and use a `left_join()` function to add the geographic information in `mapWorld` to the countries in `myeurope`. Save this new dataset as `mymap`.

```

mymap = myeurope %>%
  # recode missing variables
  mutate(country.name.en=recode(country.name.en, "Bosnia & Herzegovina"="Bosnia and Herzegovina", "Czechia"="Czech Republic", "United Kingdom"="UK")) %>%
  # combine datasets
  left_join(mapWorld, by = c("country.name.en" = "region"))

mymap %>% arrange(desc(pop_growth))

```

```

## # A tibble: 21,056 × 12
##   count...1 country year pop_u...2 pop_u...3 pop_t...4 pop_g...5 long lat group order
##   <chr> <chr> <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <int>
## 1 Luxemb... LUX 2017 543862 3.25 599449 2.95 6.12 50.1 958 59705
## 2 Luxemb... LUX 2017 543862 3.25 599449 2.95 6.11 50.1 958 59706
## 3 Luxemb... LUX 2017 543862 3.25 599449 2.95 6.11 50.0 958 59707
## 4 Luxemb... LUX 2017 543862 3.25 599449 2.95 6.14 50.0 958 59708
## 5 Luxemb... LUX 2017 543862 3.25 599449 2.95 6.20 49.9 958 59709
## 6 Luxemb... LUX 2017 543862 3.25 599449 2.95 6.26 49.9 958 59710
## 7 Luxemb... LUX 2017 543862 3.25 599449 2.95 6.32 49.8 958 59711
## 8 Luxemb... LUX 2017 543862 3.25 599449 2.95 6.44 49.8 958 59712
## 9 Luxemb... LUX 2017 543862 3.25 599449 2.95 6.49 49.8 958 59713
## 10 Luxemb... LUX 2017 543862 3.25 599449 2.95 6.49 49.8 958 59714
## # ... with 21,046 more rows, 1 more variable: subregion <chr>, and abbreviated
## # variable names 1country.name.en, 2pop_urb_total, 3pop_urb_grow, 4pop_total,
## # 5pop_growth

```

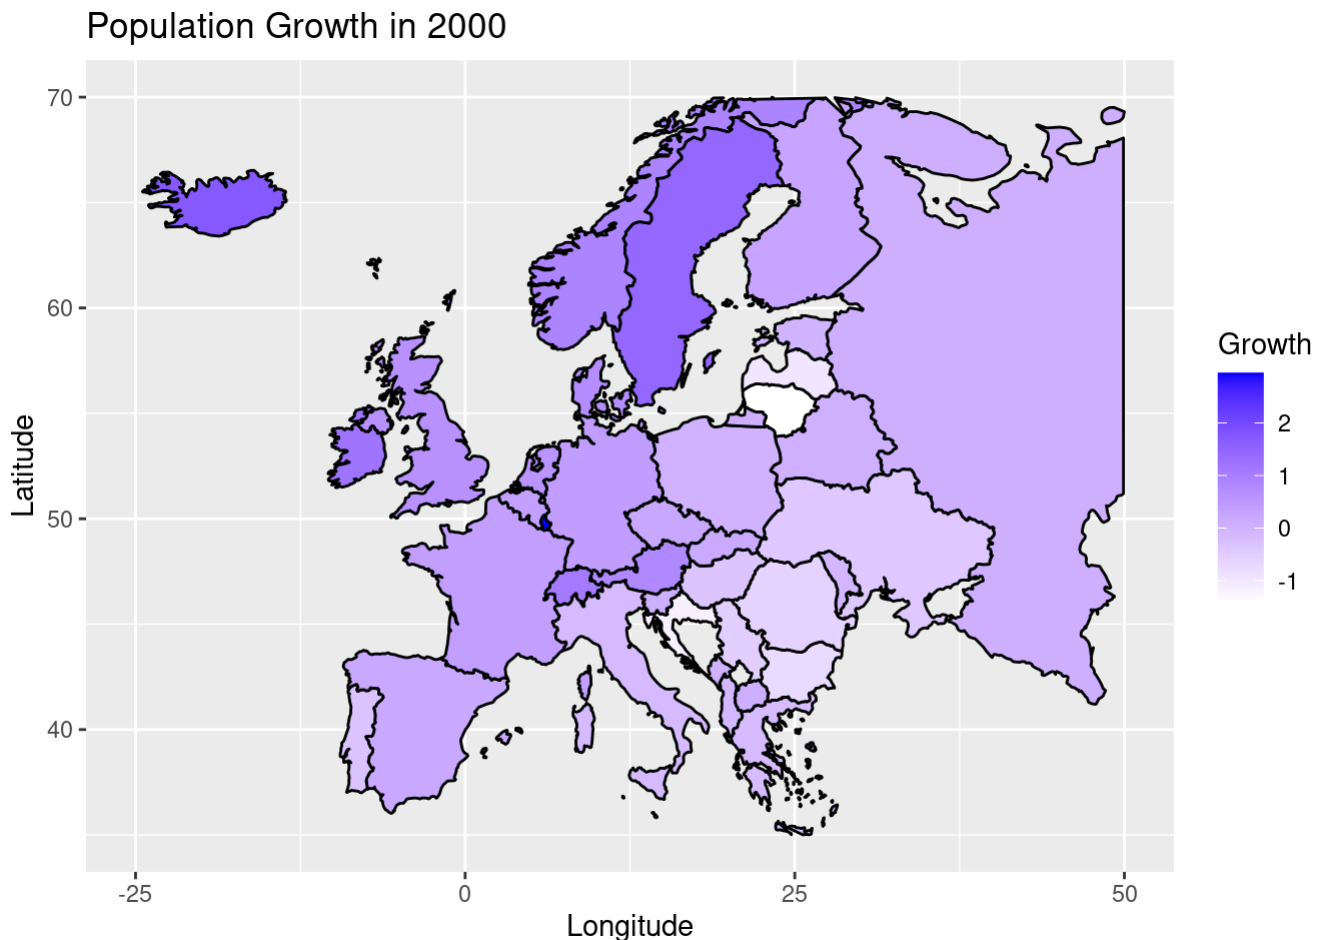
Question 12: (2 pts)

Let's visualize how population growth varies across European countries in 2017 with a map. With the package `ggmap`, use the R code provided below. Add a comment after each `#` to explain what each component of this code does. *Note: it would be a good idea to run the code piece by piece to see what each layer adds to the plot.*

```
# Paste and run the following into your console (NOT HERE): install.packages("ggmap")

# Call the ggmap package
library(ggmap)

# Build a map!
mymap %>%
  # assign data from mymap to ggplot
  ggplot(aes(x = long, y = lat, group = group, fill = pop_growth)) +
  # draw countries
  geom_polygon(colour = "black") +
  # assign fill color based on population growth
  scale_fill_gradient(low = "white", high = "blue") +
  # add title and axes labels
  labs(fill = "Growth", title = "Population Growth in 2000",
       x = "Longitude", y = "Latitude") +
  # limit spn of graph to make it look prettier
  xlim(-25,50) + ylim(35,70)
```



Which country had the highest population growth in Europe in 2017? *Hint: it's very tiny and very close to where I'm from! You can refer to this map for European geography: <https://www.wpmap.org/europe-map-hd-with-countries/> (<https://www.wpmap.org/europe-map-hd-with-countries/>)*

Luxembourg had the highest European population growth in 2017!

Formatting: (2 pts)

Comment your code, write full sentences, and knit your file!

```
##                               sysname
##                               "Linux"
##                               release
##                               "5.15.0-58-generic"
##                               version
## "#64~20.04.1-Ubuntu SMP Fri Jan 6 16:42:31 UTC 2023"
##                               nodename
##                               "educcomp01.ccbb.utexas.edu"
##                               machine
##                               "x86_64"
##                               login
##                               "unknown"
##                               user
##                               "jlh7459"
##                               effective_user
##                               "jlh7459"
```